

ISYS2120 – Data & Information Management

Week 3A: SQL DDL, and DML for data changes
(Kifer/Bernstein/Lewis - Chapter 3; Ramakrishnan/Gehrke - Chapter 3)

Based on slides from Kifer/Bernstein/Lewis (2006) “Database Systems”
and from Ramakrishnan/Gehrke (2003) “Database Management Systems”,
and including material from Fekete and Röhm.

Prof Alan Fekete



COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

WARNING

This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (**the Act**). The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice.



Schema definition

- A relational dbms knows the schema of the database it manages
 - ▶ Indeed, data cannot be stored except when the schema is already known
 - This is in contrast to some other platforms (eg JSON stores), which can accept data and then be told (or even figure out) the schema
- The schema is described in SQL statements (part of Data Definition Language, also abbreviated as DDL)
 - ▶ Give names of tables, names of columns, datatypes of columns, and also various constraints
- There are also commands that modify the schema
- Schema descriptions are stored in the system catalogue, which can be queried
 - ▶ Eg a table which gives information about the attributes of all tables

Creating Tables in SQL

■ Creation of tables:

CREATE TABLE *name* (*list-of-columns-with-types*)

- ▶ Example: Create the Students relation.

Observe that the datatype (domain) of each field is specified, and enforced by the DBMS whenever tuples are added or modified.

```
CREATE TABLE Instructor (lname VARCHAR(20) ,  
                             fname  VARCHAR(20) ,  
                             salary  INTEGER,  
                             birth   DATE ,  
                             hired   DATE ) ;
```

- ▶ Remember that table and column names are SQL identifiers (and case-insensitive)
- ▶ Remember not to use any keyword as an identifier (eg, don't try to have a column called `date`)
- ▶ After performing this, the table exists, but it starts empty (contains no rows yet)



Base Datatypes of SQL

Base Datatypes	Description	Example Values
SMALLINT INTEGER BIGINT	Integer values	1704, 4070
DECIMAL(p,q) NUMERIC(p,q)	Fixed-point numbers with precision p and q decimal places	1003.44, 160139.9
FLOAT(p) REAL DOUBLE PRECISION	floating point numbers with precision p	1.5E-4, 10E20
CHAR(q) VARCHAR(q) CLOB(q)	alphanumeric character string types of fixed size q respectively of variable length of up to q chars	'The quick brown fox jumps...', 'ISYS2120'
BLOB(r)	binary string of size r	B'01101', X'9E'
DATE	date	DATE '1997-06-19', DATE '2001-08-23'
TIME	time	TIME '20:30:45', TIME '00:15:30'
TIMESTAMP	timestamp	TIMESTAMP '2002-08-23 14:15:00'
INTERVAL	time interval	INTERVAL '11:15' HOUR TO MINUTE

(cf. Türker, ORDBMS 2004/2005)



Create Table Example

<i>Student</i>	
sid	name

<i>Enrolled</i>		
sid	ucode	grade

<i>UnitOfStudy</i>		
ucode	title	credit_pts

CREATE TABLE Student (-- Note: layout on lines is not significant!
 sid **INTEGER**,
 name **VARCHAR(20)**
);
CREATE TABLE UnitOfStudy
(
 ucode **CHAR(8)**,
 title **VARCHAR(30)**,
 creditPoints **INTEGER**
);
CREATE TABLE Enrolled (
 sid **INTEGER**, ucode **CHAR(8)**, grade **INTEGER**
);



The Special NULL 'Value'

- SQL-based RDBMS allows a special entry **NULL** in a column to represent facts that are not relevant, or not yet known
 - This can occur in a column of any datatype
 - This is different than the VARCHAR 'Null'
- ▶ Eg a new employee has not yet been allocated to a department
- ▶ Eg salary, hired may not be meaningful for adjunct lecturers
- ▶ Eg INSTRUCTOR table in a university

Iname	fname	salary	birth	hired
Jones	Peter	35000	1970	1998
Smith	Susan	null	1983	null
Smith	Alan	35000	1975	2000

Evaluation of use of NULL

■ Advantage:

NULL is useful because using an ordinary value with special meaning does not always work

- ▶ Eg if salary=-1 is used for “unknown” in the previous example, then averages won’t be sensible

■ Disadvantage:

NULL causes complications in the definition of many operations (eg is $27 < \text{NULL}$? Or $27 > \text{NULL}$?)

SQL Schemas & Table Modifications

- Database Servers typically shared by multiple users
 - ▶ We want separate schemas per user (naming: `schema.tablename`)
 - ▶ **CREATE SCHEMA** ... command
 - E.g. http://download.oracle.com/docs/cd/B19306_01/server.102/b14200/statements_6014.htm
or <http://www.postgresql.org/docs/8.3/static/sql-createschema.html>
 - ▶ If not provided, automatic schema by user name (cf. Oracle)
- Several base data types available in SQL
 - ▶ E.g. `INTEGER`, `REAL`, `CHAR`, `VARCHAR`, `DATE`, ...
 - ▶ but each system has its specialities such as specific BLOB types or value range restrictions
 - E.g. Oracle calls a string for historical reasons `VARCHAR2`
 - ▶ cf. online documentation
- Existing schemas can be changed
 - ALTER TABLE** *name* **ADD COLUMN** ... | **ADD CONSTRAINT**... | ...
 - ▶ Huge variety of vendor-specific options; cf. online documentation



Integrity Constraints

- **Integrity Constraint (IC):** a condition that must be true for *any* instance of the database; e.g., domain constraints, uniqueness constraints and more complicated ones too.
 - ▶ ICs can be declared in the schema
 - They are specified when schema is defined.
 - All declared ICs are checked whenever relations are modified.
- A *legal* instance of a relation is one that satisfies all specified ICs.
 - ▶ If ICs are declared, DBMS will not allow illegal instances.
- When the DBMS checks ICs, stored data is more faithful to real-world meaning.
 - ▶ Can avoid some data entry errors, too!



Non-Null Columns

- One domain constraint is to insist that no value in a given column can be null
 - ▶ The value can't be unknown; The concept can't be inapplicable
- In SQL, append NOT NULL to the field declaration

CREATE TABLE Instructor

```
( lname VARCHAR(20) NOT NULL,  
  fname VARCHAR(20) NOT NULL,  
  salary INTEGER,  
  birth DATE NOT NULL,  
  hired DATE );
```



Primary Key

- Recall that many entities have identifiers, which can be used to distinguish one entity from another in the entity set, even when many attributes are the same
 - ▶ Often, an identifier is completely artificial, and the value has no meaning at all
 - ▶ Sometimes, the identifier is a sequence number
- In a relational schema, we can declare that a column is a Primary Key (often abbreviated as PK in textbooks or diagrams)
 - ▶ This allows it to be used as identifier in other tables, to connect information
 - ▶ It will enforce that every row has a different value for this column
 - ▶ It will enforce that no row has NULL as value for this column

SQL ways to indicate a PK

- Append to end of column definition

```
CREATE TABLE Student (  
    sid    INTEGER PRIMARY  
    KEY,  
    name VARCHAR(20)  
);
```

- Extra clause in table creation

```
CREATE TABLE Student (  
    sid    INTEGER,  
    name VARCHAR(20),  
    PRIMARY KEY (sid)  
);
```

- Named constraint in table creation

```
CREATE TABLE Student (  
    sid    INTEGER,  
    name VARCHAR(20),  
    CONSTRAINT Student_PK  
        PRIMARY KEY (sid)  
);
```

- Separate named constraint added later

```
CREATE TABLE Student (  
    sid    INTEGER,  
    name VARCHAR(20)  
);  
ALTER TABLE Student (  
    ADD CONSTRAINT Student_PK  
    PRIMARY KEY (sid)  
);
```

Composite PK

- In some tables, no single column is an identifier, but a combination of columns may be able to distinguish the rows
- Eg, sometimes, each manufacturer has its own set of serial numbers,
 - ▶ it is possible to have two rows with the same serialno, when they have different manufacturer
 - ▶ And it is possible to have two rows with the same manufacturer, when they have different serialno
 - ▶ But the combination (manufacturer, serialno) is suitable as an identifier
 - ▶ This is called a composite primary key

■ In SQL

```
CREATE TABLE Product (  
  manufacturer VARCHAR(20),  
  serialno INTEGER,  
  PRIMARY KEY (manufacturer, serialno)
```

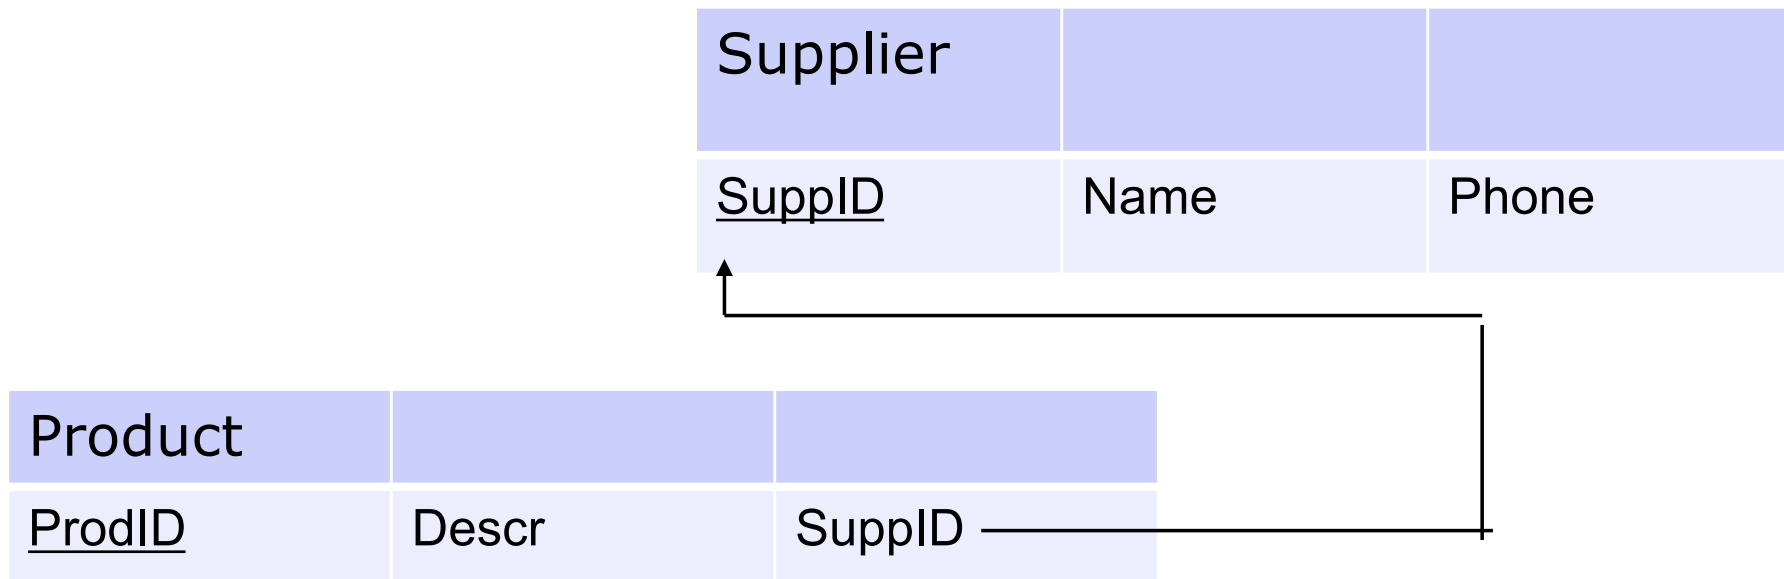


Uniqueness

- A table can have only one primary key declared (may be a single column, or a combination of columns)
- Sometimes, there are other columns or combinations that are necessarily different among rows
 - ▶ We call any combination of columns which can distinguish the rows, as a candidate key
- SQL allows multiple constraints on a table, which say a column or combination is UNIQUE

Foreign Key

- When we use values to connect information across tables, we expect that there will be something with that value in the referenced table
- We can make this explicit with a FOREIGN KEY constraint (often abbreviated FK in text or diagrams)
- Recall the relational schema diagram



Foreign Key Constraint in SQL

- Several ways to express this, among them

```
CREATE TABLE Product ( ProdID INTEGER,  
    descr VARCHAR(10),  
    SuppID INTEGER REFERENCES Supplier(SuppID));
```

Or

```
CREATE TABLE Product ( ProdID INTEGER,  
    descr VARCHAR(10),  
    SuppID INTEGER REFERENCES Supplier);
```

Or -- When no column is mentioned for referenced table, implicit reference to its primary key

```
CREATE TABLE Product ( ProdID INTEGER,  
    descr VARCHAR(10),  
    SuppID INTEGER,  
    FOREIGN KEY (SuppID) REFERENCES Supplier(SuppID));
```

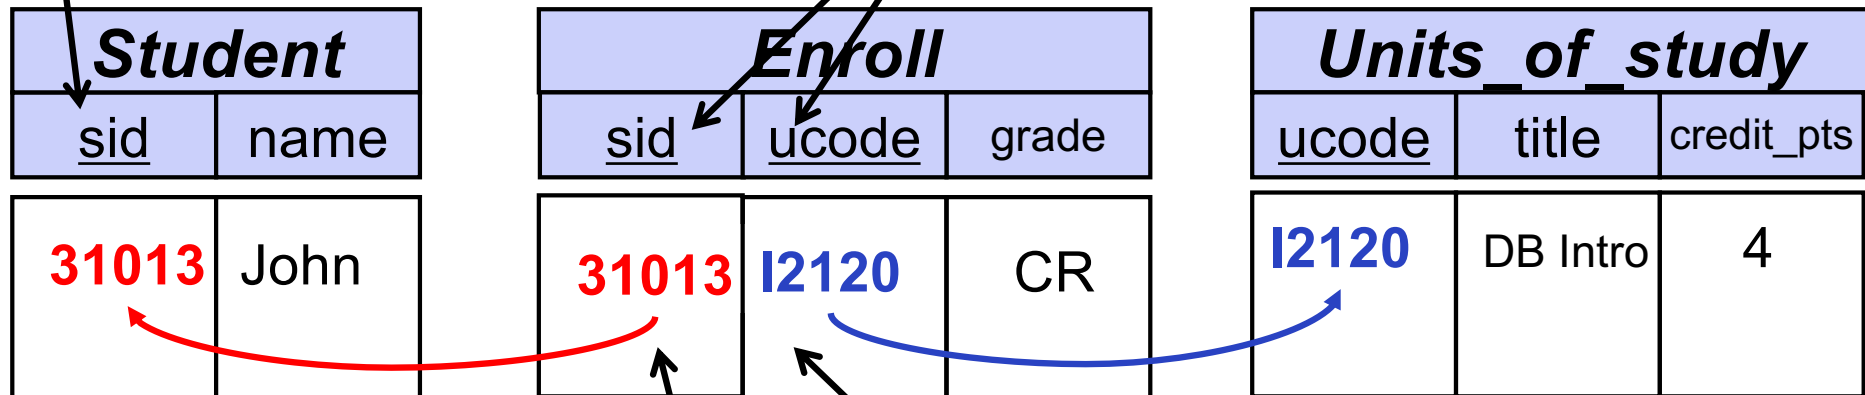
Or

```
CREATE TABLE Product ( ProdID INTEGER,  
    descr VARCHAR(10),  
    SuppID INTEGER,  
    CONSTRAINT Product_FK FOREIGN KEY (SuppID) REFERENCES  
    Supplier(SuppID));
```

Summary of terms

Primary key identifies each tuple of a relation.

Composite Primary Key consisting of more than one attribute.



Foreign key is a (set of) attribute(s) in one relation that 'refers' to a tuple in another relation by giving the value of its key (like a 'logical pointer'). Another foreign key

Create Table Example with PKs/FKs

<i>Student</i>	
<u>sid</u>	name

<i>Enrolled</i>		
<u>sid</u>	<u>ucode</u>	grade

<i>Unit_of_Study</i>		
<u>ucode</u>	title	credit_pts

```
CREATE TABLE Student ( sid INTEGER, ... ,  
    CONSTRAINT Student_PK PRIMARY KEY (sid)  
);
```

```
CREATE TABLE UoS ( ucode CHAR(8), ... ,  
    CONSTRAINT UoS_PK PRIMARY KEY (ucode)  
);
```

```
CREATE TABLE Enrolled ( sid INTEGER, ucode CHAR(8), grade CHAR(2),  
    CONSTRAINT Enrolled_FK1 FOREIGN KEY (sid) REFERENCES Student,  
    CONSTRAINT Enrolled_FK2 FOREIGN KEY (ucode) REFERENCES UoS,  
    CONSTRAINT Enrolled_PK PRIMARY KEY (sid,ucode)  
);
```



Choosing the Correct Key Constraints

- Careful: Used carelessly, an IC can prevent the storage of database instances that arise in practice!
- Example:
Attempt to model that a student can get only a single grade per course.

<pre>CREATE TABLE Enrolled (sid INTEGER, cid CHAR(8), grade CHAR(2), PRIMARY KEY (sid,cid))</pre>	vs.	<pre>CREATE TABLE Enrolled (sid INTEGER, cid CHAR(8), grade CHAR(2), PRIMARY KEY (sid, cid), UNIQUE (sid, grade))</pre>
--	-----	--

■ “For a given student and course, there is a single grade; the same grade can be achieved by several students in a course.”

■ “For a given student and course, there is a single grade; but a student can achieve a *certain grade only once*.”



Brainteaser

- Given the following example a table

```
CREATE TABLE Test (  
    a INTEGER,  
    b INTEGER UNIQUE,  
    PRIMARY KEY (a,b)  
);
```

- Would the following be a legal database instance?

```
{ (1, 1) ,  
  (1, 2) ,  
  (1, 3) ,  
  (2, 1) ,  
  (2, 4) }
```



Keys and NULLs

■ PRIMARY KEY

- ▶ Must be unique and do not allow NULL values

■ UNIQUE (candidate key)

- ▶ Possibly many *candidate keys* (specified using UNIQUE)
- ▶ According to the ANSI standards SQL:92, SQL:1999, and SQL:2003, a UNIQUE constraint should disallow duplicate non-NULL values, but allow multiple NULL values.
- ▶ Many DBMS (e.g. Oracle or SQL Server) implement only a crippled version of this, allowing a single NULL but disallowing multiple NULL values....

■ FOREIGN KEY

- ▶ By default allows nulls
- ▶ If there must be a parent tuple, then must combine with NOT NULL constraint



Foreign Keys & Referential Integrity

■ Referential Integrity:

for each tuple in the referring relation whose foreign key value is **X**, there must be a tuple in the referred relation with a candidate key that also has value **X**

▶ e.g. *sid* is a foreign key referring to Student:

Enrolled(*sid*: integer, ucode: string, grade: string)

■ When all foreign key constraints are enforced, referential integrity is achieved, i.e., no dangling references

Q: Can you name a data model w/o referential integrity?



How to enforce Referential Integrity

- Consider Student and Enrolled;
sid in Enrolled is a foreign key that references Student.
- What should be done if an Enrolled tuple with a non-existent student *sid* is inserted? (*Reject it!*)
- What should be done if a Student tuple is deleted? Choices:
 - ▶ Also delete all Enrolled tuples that refer to it.
 - ▶ Disallow deletion of a Student tuple that is referred to.
 - ▶ Set *sid* in Enrolled tuples that refer to it to a *default sid*.
 - ▶ (In SQL, also: Set *sid* in Enrolled tuples that refer to it to a special value *null*, denoting '*unknown*' or '*inapplicable*'.)
- Similar if primary key of Student tuple is updated.



Referential Integrity in SQL

- SQL/92, SQL:1999 and SQL:2003 support all 4 options on deletes and updates.
 - ▶ Default is **NO ACTION** (*delete/update is rejected*)
 - ▶ **CASCADE** (also delete all tuples that refer to deleted tuple)
 - ▶ **SET NULL / SET DEFAULT** (sets foreign key value of referencing tuple)

```
CREATE TABLE UnitOfStudy
( uosCode    CHAR(8) ,
  title      VARCHAR(80) ,
  credit_pts  INTEGER ,
  taughtBy   INTEGER DEFAULT 1 ,
  PRIMARY KEY (uosCode) ,
  FOREIGN KEY (taughtBy)
  REFERENCES Professor
  ON UPDATE CASCADE
  ON DELETE SET DEFAULT )
```

UnitOfStudy

<u>uosCode</u>	title	credit_pts	taughtBy
----------------	-------	------------	----------

Professor

<u>empid</u>	name
--------------	------



CHECK Constraints

- In SQL the domain of a column is given as a simple datatype
- Sometimes, we know more about the valid values that can be there
 - ▶ Eg Salary should be positive
 - ▶ Eg UoSCode should be 4 letters then 4 digits
- We can define a Boolean property that is required to always be true for every row

```
CREATE TABLE Employee
( empID    INTEGER,
  ...
  salary   INTEGER CHECK (salary > 0), ... );
```

- If the property involves several columns, it needs to be a separate clause in the SQL, not just appended to the column declaration

Changing schema in SQL

■ Deletion of a tables:

DROP TABLE *name*

- ▶ the schema information and the tuples in the table, are all removed.
- ▶ Example: Destroy the Instructor relation

DROP TABLE Instructor

■ Existing schemas can be changed

ALTER TABLE *name* **ADD COLUMN** ... | **ADD CONSTRAINT**... | ...

- ▶ Huge variety of vendor-specific options; cf. online documentation



Modifying Instances using SQL DML

■ Insertion of new data into a table / relation

▶ Syntax:

INSERT INTO *table* [(*"list-of-columns"*)] **VALUES** (*" list-of-expression "*)

▶ Example:

```
INSERT INTO Student (sid, name) VALUES (12345678, 'Smith')
```

■ Updating of tuples in a table / relation

▶ Syntax:

UPDATE *table* **SET** *column* = *"expression"* {*","column* = *"expression"*}
[**WHERE** *search_condition*]

▶ Example: UPDATE Student

```
    SET address = '4711 Water Street'  
WHERE sid = 123456789
```

■ Deleting of tuples from a table / relation

▶ Syntax:

DELETE FROM *table* [**WHERE** *search_condition*]

▶ Example:

```
DELETE FROM Student WHERE name = 'Smith'
```

WHERE clause
as in SELECT statement



Recall terminology

■ Informal Definition:

A **relation** is a named, two-dimensional table of data

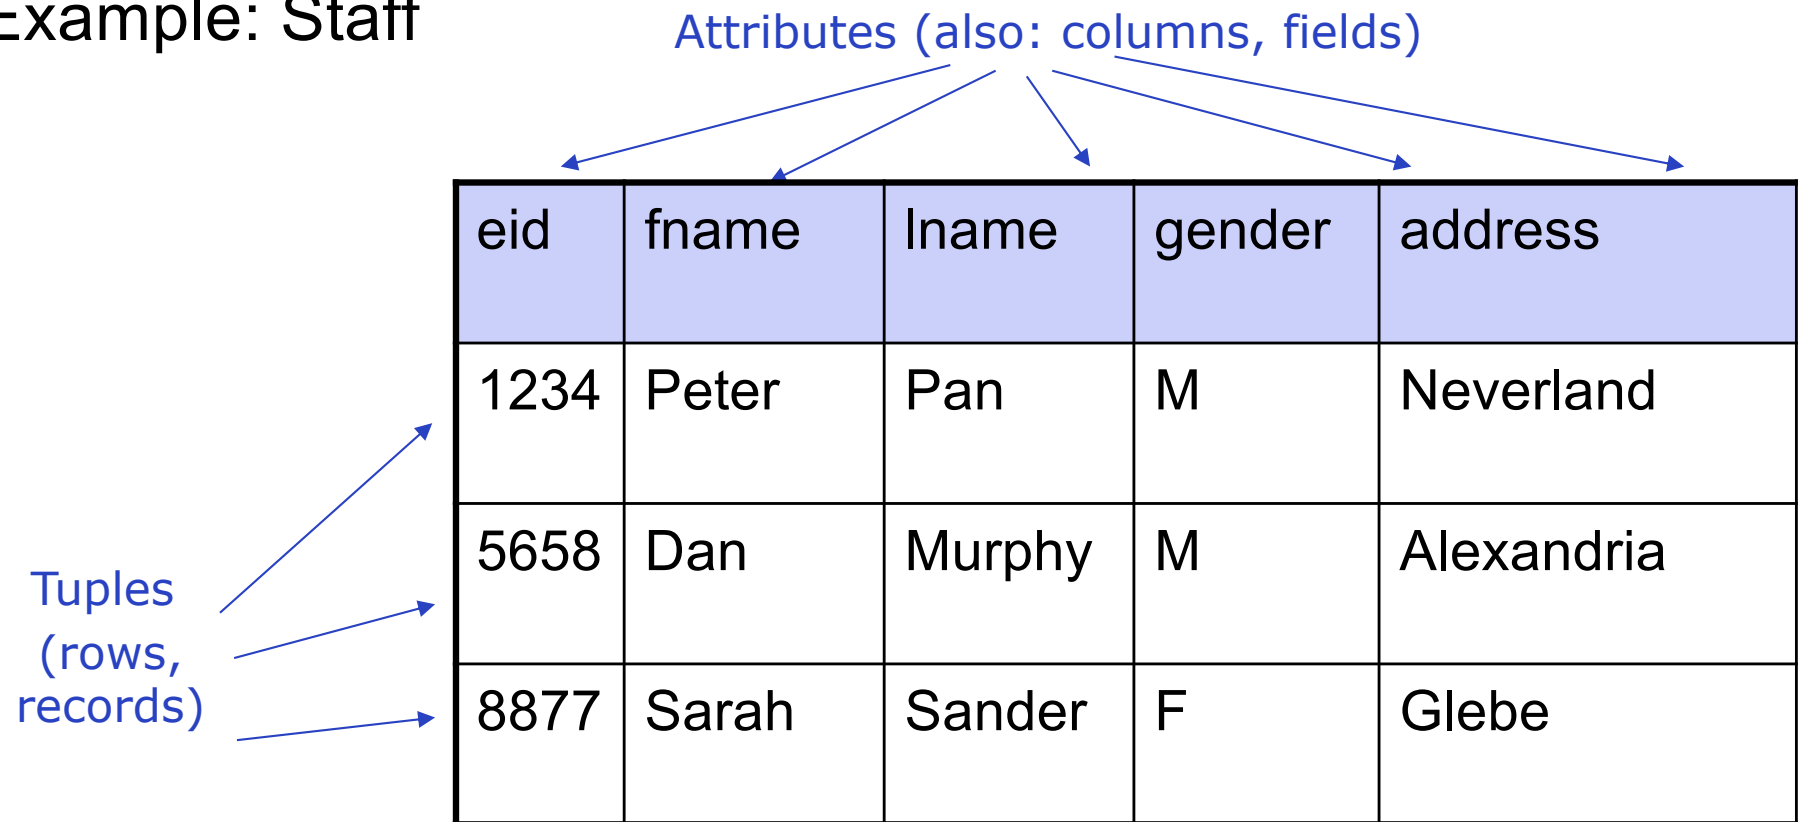
- ▶ Table consists of rows (record) and columns (attribute or field)

■ Example: Staff

Attributes (also: columns, fields)

eid	fname	lname	gender	address
1234	Peter	Pan	M	Neverland
5658	Dan	Murphy	M	Alexandria
8877	Sarah	Sander	F	Glebe

Tuples (rows, records)



Maths Definition of a Relation

- A Relation is a mathematical concept from discrete maths based on the ideas of sets.

- ▶ **Relation R**

Given sets D_1, D_2, \dots, D_n , a relation R is a subset of $D_1 \times D_2 \times \dots \times D_n$

Thus, a relation is a set of n -tuples (a_1, a_2, \dots, a_n) where $a_i \in D_i$

- Example:

If

$studentid = \{12345678, 23456789, 345354345, 44455666,$
 $etc\}$

$name = \{Jones, Smith, Kerry, Lindsay, etc\}$

$date = \{1985-11-09, 1984-07-15, 1984-12-01, 1986-01-01,$
 $etc\}$

then

$R = \{ (12345678, Jones, 1984-07-15),$
 $(345354345, Lindsay, 1986-01-01),$
 $(44455666, Kerry, 1985-11-09),$
 $(23456789, Kerry, 1994-07-15) \}$

is a relation over $studentid \times name \times date$



Theory vs. Technology

- A relational DBMS supports data in a form close to, but not exactly, matching the mathematical relation
 - ▶ RDBMS allows null values for unknown information
 - ▶ RDBMS gives a name for each position among the columns
 - ▶ RDBMS insists that datatype for any column is from a limited variety, all simple
 - ▶ RDBMS considers the columns as being arranged in an order
 - ▶ RDBMS considers the rows as being arranged in an order
 - ▶ RDBMS allows duplicate rows



References

- Kifer/Bernstein/Lewis (2nd edition)
 - ▶ Chapter 3
- Ramakrishnan/Gehrke (3rd edition - the 'Cow' book)
 - ▶ Chapter 3.1-3.4 and 3.6-3.7, plus Chapter 1.5
- Ullman/Widom (3rd edition)
 - ▶ Chapter 2.1 - 2.3, Section 7.1 and Chapter 8.1-8.2
 - ▶ *views and foreign keys come later, instead relational algebra is introduced very early on; also briefly compares RDM with XML,*
- Silberschatz/Korth/Sudarshan (5th edition - 'sailing boat')
 - ▶ Chapter 2.1 - 2.2; Chapter 3.1-3.2 and 3.9
 - ▶ *starts with relational algebra early on which we do later*
- Elmasri/Navathe (5th edition)
 - ▶ Chapter 2.1-2.3; Chapter 5; Section 8.8
 - ▶ *talks first more about system architectures and conceptual modeling*