

# **ISYS2120: Data & Information Management**

## **Week 8: Database Application Development**

**Alan Fekete**

Based on material from Roehm, Khushi

Cf. Kifer/Bernstein/Lewis – Chapter 8


Ramakrishnan/Gehrke – Chapter 6

Silberschatz/Korth/Sudarshan – Chapter 9

Ullman/Widom – Chapter 9



THE UNIVERSITY OF  
**SYDNEY**



## **COMMONWEALTH OF AUSTRALIA**

### **Copyright Regulations 1969**

#### **WARNING**

This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (**the Act**). The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

**Do not remove this notice.**

# Database-backed software

- Most of the software used for life and business needs to use data that is found in a database
  - ▶ Often, user activities will also modify the data in the database
- A few examples
  - ▶ E-commerce (inventory, purchase, order status, order history)
  - ▶ Entertainment (catalogue, preferences, history)
  - ▶ Social media (posts, community connections)
  - ▶ Transport (routes, timetables, current status)
- The software that provides these functionalities, needs to access one or more dbms

# Software and queries

- Previously in isys2120, we covered how to write SQL to extract and modify data
- A human typed the query into a query window and ran it against the database, and observed the output
- End-users can't be expected to know SQL, nor usually to have accounts and access on all the dbms needed for modern life
- Instead, the user invokes some software, and the software submits queries to dbms, gets result table, and displays information to the user
- This lecture is about the structure of that software

# Data-intensive Systems

- Three types of functionality (often placed in separate layers of code):

## **Presentation Logic**

- Input – keyboard/mouse/gestures
- Output – monitor/printer/screen

**GUI Interface**

## **Processing Logic**

- Business rules
- I/O processing

**Procedures, functions, programs**

## **Data Management**

(Storage Logic)

- data storage and retrieval

**DBMS activities**

- The system architecture determines whether these three components reside on a single computing system (1-tier) or whether they are distributed across several tiers

# Presentation layer

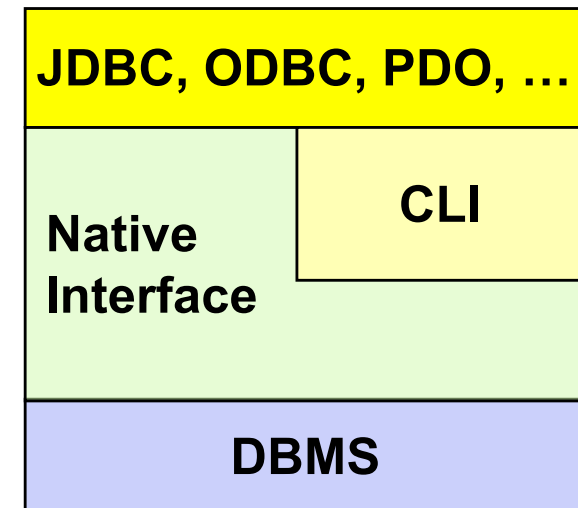
- Often, web browser is used as GUI for end-user
  - ▶ Available on all devices!
- Application needs to have code that lays out the web pages and provides navigation between them
- Mobile devices may provide an app that directly works with the gestures etc of the device
  - ▶ And is targeted for the small screen size
- Lab08 and Asst3 will work with web interface
  - ▶ Flask library of Python, to construct a web-server that also runs business logic and data management

# SQL in Application Code

- SQL commands can be called from within a *host language* (such as Python or Java) program.
  - ▶ Must include a statement to *connect* to the right database.
  - ▶ SQL statements can refer to *host variables* (including special variables used to return status).
  
- Two main integration approaches:
  - ▶ **Statement-level interface (SLI)**
    - Embed SQL in the host language (Embedded SQL in C, SQLJ)
    - Application program is a mixture of host language statements and SQL statements and directives
    - Compiler must deal with both aspects
  - ▶ **Call-level interface (CLI)**
    - Create special API to call SQL commands (JDBC, ODBC, Python, ...)
    - SQL statements are passed as arguments to host language (library) procedures / APIs

# Call-level Interfaces and Database APIs

- Program can invoke methods/procedures in a library with database calls (API)
  - ▶ Pass SQL strings from language, present result sets in language-friendly way
  - ▶ Supposedly DBMS-neutral
    - a “driver” executes the calls and translates them into DBMS-specific code
    - database can be across a network
- Several Variants
  - ▶ **SQL/CLI**: “SQL Call-Level-Interface”
    - Part of the SQL-92 standard;
    - “The assembler under the APIs”
  - ▶ **ODBC**: “Open DataBase Connectivity”
    - Side-branch of early version of SQL/CLI
    - Enhanced to: **OLE/db**, and further **ADO.NET**
  - ▶ **JDBC**: “Java DataBase Connectivity”
    - Java standard
  - ▶ **PDO**
    - Persistency standard for PHP Data Objects





# PYTHON DB-API2

a Call-Level API Example



# Python

- Python features extensive standard library (modules)
  - ▶ Special functionality supported by variety of optional 3<sup>rd</sup>-party modules
  - ▶ For database connectivity, several database-specific python modules
    - e.g. psycopg (PostgreSQL) or cx\_oracle (Oracle)
    - **<http://initd.org/psycopg/docs/>**
  - ▶ For dynamic websites:
    - several framework available; we will use Flask..
    - Allows to define template pages with embedded python code

# Python Database API Specification (DB-API)

- DB-API 2.0 was released April 1999
- Defines common functions and API for access modules to different database systems
  - ▶ Module API; Connection and Cursor interface definitions
- Works as a generic as a **database abstraction layer**
  - ▶ Generic driver model to connect to different database engines via the same API
- URLs:
  - <https://wiki.python.org/moin/DatabaseProgramming>
  - <https://wiki.python.org/moin/UsingDbApiWithPostgres>
  - <http://initd.org/psycopg/docs>
  - [http://www.tutorialspoint.com/postgresql/postgresql\\_python.htm](http://www.tutorialspoint.com/postgresql/postgresql_python.htm)
  - <https://www.python.org/dev/peps/pep-0249/>

# Python DB-API Example

```
import psycopg2

try:
    # connect to the database
    conn = psycopg2.connect(database="postgres",user="test",password="secret")

    # prepare to query the database
    curs = conn.cursor()

    # execute a parameterised query
    unit_of_study = "ISYS2120"
    curs.execute("""SELECT name
                    FROM Student NATURAL JOIN Enrolled
                    WHERE uos_code = %(uos)s""", {'uos': unit_of_study} )

    # loop through the resultset
    for result in curs:
        print (" student: " + result[0])

    # clean up
    curs.close()
    conn.close()

except Exception as e: # error handling
    print("SQL error: unable to connect to database or execute query")
    print(e)
```



# Core tasks with SQL Interfaces

(1) Establishing a database connection

(2) Static vs. Dynamic SQL

(3) Parameterized SQL and mapping of domain types to data types of host

- ▶ Concept of *host variable*
- ▶ How to treat *NULL* values?

(4) Impedance Mismatch:

- ▶ SQL operates on sets of tuples
- ▶ Host languages like C do not support a set-of-records abstraction, but only a one-value-at-a-time semantic
- ▶ Solution: *Cursor Concept*  
Iteration mechanism (loop) for processing a set of tuples

(5) Error handling



# (1) DB Connections from Python

- Session with PostgreSQL started by creating a connection

- Two Variants:

- ▶ Connect with keyword arguments

```
conn = psycopg2.connect(host='...', database='...', user='X', password='...')
```

- ▶ Connect with a Data Source Name (*DSN*) string of the form

**"host=X dbname=Y user=U password=P"**

For example for PostgreSQL:

```
conn = psycopg2.connect(  
    "host=postgres.usyd.edu.au dbname=unldb user=U password=secret")
```

*connectionParameters*                      *db login*

Details: <http://initd.org/psycpg/docs/module.html#psycpg2.connect>

# Python Database Connection Modules

## ■ Python support for variety of DBMSs

- ▶ MySQL (module: MySQLdb)
- ▶ PostgreSQL (module: psycopg2)
- ▶ Oracle (module: cx\_oracle)
- ▶ IBM DB2 (module: ibm\_db)
- ▶ SQL Server (module: pymssql)
- ▶ sqlite (module: sqlite3)
- ▶ ...
- ▶ DSN syntax and additional DB parameters vary for each driver
  - Check manuals...

### **Note:**

db modules need to be installed first as part of the Python installation...

## ■ Example for Oracle:

```
dsnStr = cx_oracle.makedsn("oracle10g.it.usyd.edu.au",1521,"ORCL")
conn = cx_oracle.connect(user="myuser",password="mypass",dsn=dsnStr)
```

# psycopg Connection Simple Example

```
import psycopg2

# connect to the database
try:
    conn = psycopg2.connect(database='foo', user='dbuser', password='pwd')
except:
    print("unable to connect to database")

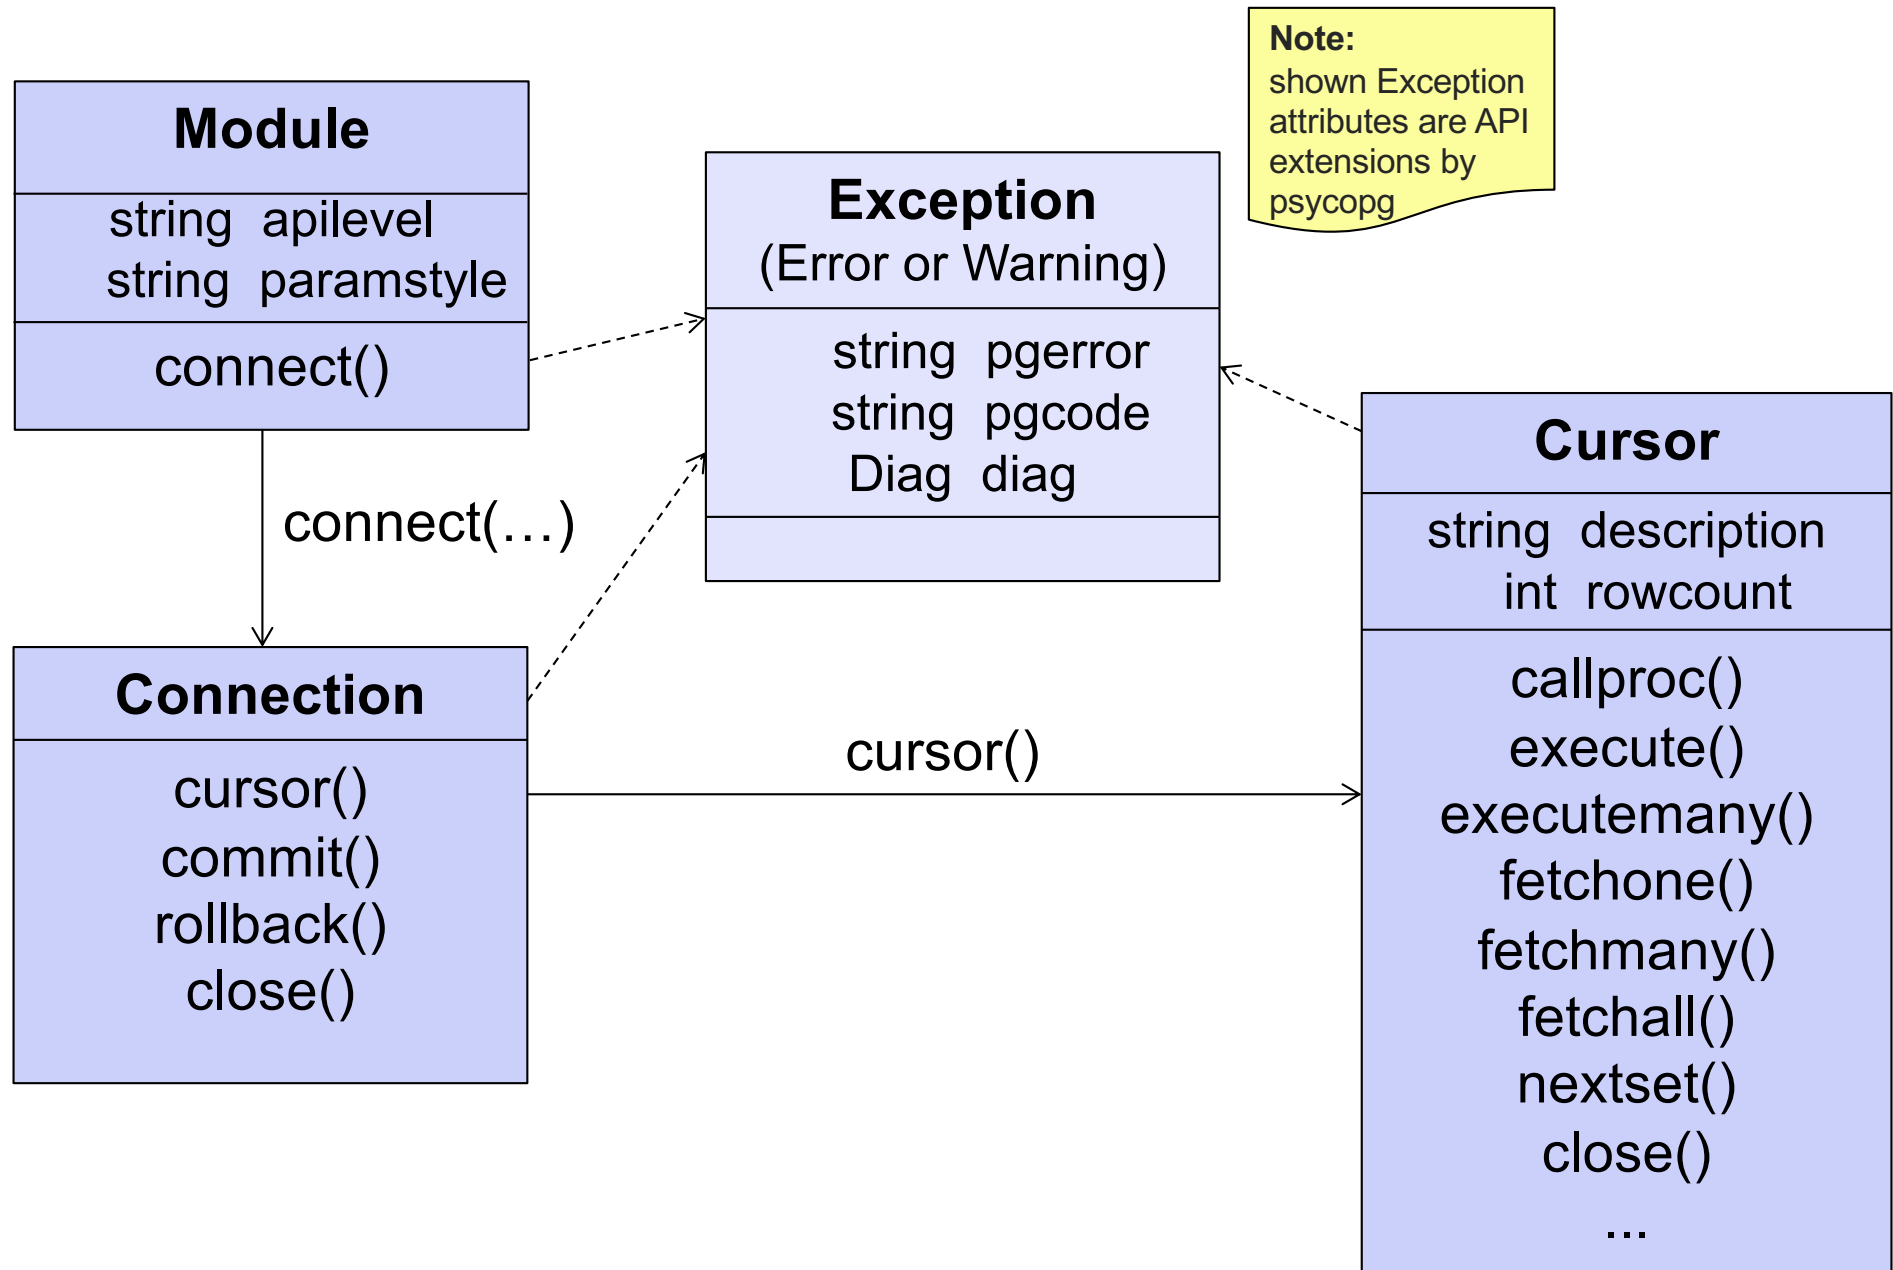
# query database
curs = conn.cursor()
try:
    curs.execute("SELECT name FROM Student WHERE studID=4711")
except:
    print("unable to execute query")

... Do Actual Work ...

# cleanup
curs.close()
conn.close()
```



# Python DB-API 2.0 Objects



# Python Database API Interfaces

- Connection Management
  - ▶ `psycopg2.connect()` connects to a database
  - ▶ `conn.cursor()` creates a cursor object for query execution
- Start SQL statements
  - ▶ `execute()` for static SQL, and also parameterized SQL queries
  - ▶ `callproc()` for executing a stored procedure including parameters
- Result retrieval
  - ▶ `fetchone()` retrieves next row of a result or **None** when no more data
  - ▶ `fetchall()` retrieves the whole (remaining) result set, and returns it as a list of tuples
- Transaction control
  - ▶ `conn.commit()` successfully finishes (commits) current transaction
  - ▶ `conn.rollback()` aborts current transaction
- Error Handling
  - ▶ Via standard exception handling of Python

# Side Note on DB Connections

- Establishing a database connection takes some time...
  - ▶ Network communication, memory allocation, dbs authorization
- So do this only once in your program
  - ▶ ... but **not** for individual SQL queries
- Modern, multi-threaded applications will typically want to have a pool of connections that are re-used
  - ▶ Might be handled by your runtime library (that's what happens in Python)
  - ▶ But for, e.g., Java programs better be mindful of connection costs!

## (2) Static vs. Dynamic SQL

### ■ SQL constructs in an application can take two forms:

#### ▶ Static SQL statements:

Useful when SQL query is fully known at compile time

- no parameters are allowed in the query string
- only useful in context of compiled languages such as C

#### ▶ Dynamic SQL statements:

Application constructs SQL statements at *run time* as values of host language variables that are manipulated by directives.

- Challenge: Python is not a compiled language;  
everything in Python/psycopg2 is by definition dynamic SQL...
- This means we have to be careful on how we construct any query and in particular how parameters are passed to the database

# DB-API: Executing SQL Statements

- Three different ways of executing SQL statements:
  - ▶ **`cursor.execute(sql)`** semi-static SQL statements
  - ▶ **`cursor.execute(sql,params)`** parameterized SQL statements
  - ▶ **`cursor.callproc(call,args)`** invoke a stored procedure
  - ▶ **`cursor.executemany(sql,seq_of_params)`** repeatedly executes parameterized SQL statements
  
- In DB-API 2.0,
  - Need to create new cursor and re-issue SQL statement each time when parameters change – or if possible use **`executemany()`**
  - Some other APIs offer “prepared statements” – parsed and optimized once in the dbms, then re-executed over and over with different parameters

# Python DB-API with fixed SQL

- Simplest way to execute some static SQL query:

```
import psycopg2

try:
    # connect to the database
    conn = psycopg2.connect(database='foo', user='dbuser', password='pwd')

    # query database
    curs = conn.cursor()
    curs.execute("SELECT name FROM Student WHERE studID=4711")
    result = curs.fetchone()
    print(result)

    # cleanup
    curs.close()
    conn.close()

except:
    print("unable to connect to db or to execute query")
```

# DB-API: Batch Insert Example

## ■ Example: executing batch INSERT statements

```
import psycopg2

try:
    # connect to the database
    conn = psycopg2.connect(database='foo', user='dbuser', password='pwd')

    # prepare list of insert values (3 students enrolling in INFO2120)
    params = [(4711, 'INFO2120'), (4712, 'INFO2120'), (4713, 'INFO2120')]

    # execute INSERT statement batch
    curs = conn.cursor()
    curs.executemany("INSERT INTO Enrolled VALUES (%s,%s)", params)
    conn.commit() # cf. next week on transactions

    # cleanup
    curs.close()
    conn.close()

except:
    print("unable to connect to db or to execute query")
```

# WARNING:

## Never Query by "String-Concatenation"

Never, never, NEVER ever use Python string concatenation (+) or string parameter interpolation (%) to place variables into a SQL query string!

=> otherwise your program is vulnerable to SQL Injection attacks

```
query = """SELECT  E.studId  FROM  Enrolled E
           WHERE  E.uosCode = """ + uosCode + \
           " AND E.semester = """ + semester
```

*string concatenation*

```
cursor.execute( query )
```

- *simple approach to construct a variable query*
- *concatenates query from different string-parts*
- *executes the constructed SQL string directly in DBMS*
- **Warning:** *prone to errors and even hacking when input strings allowed to be entered by a user*



# SQL Code Injection Vulnerability

## ■ SQL-Injection

to infiltrate a SQL database with own SQL commands.

- ▶ Can be used to execute SQL statements with elevated privileges or to impersonate another user.

## ■ Without direct database connection (e.g. web application)

- ▶ Injecting SQL via un-checked user input.
- ▶ Exploiting buffer overflows.
  - Oracle standard packages have many buffer overflows.
- ▶ Output on attacker's screen.

## ■ With a direct database connection

- ▶ SQL Injection in built-in or user-defined procedures.
- ▶ Buffer overflows in built-in or user-defined procedures.
  - Risk when a procedure is not defined with the AUTHID CURRENT\_USER keyword (executes with the privileges of the owner)

# DB-API: Parameterized Queries

- Two (safe) approaches for passing query parameters:  
(because `execute()` will do any necessary escaping / conversions for parameter markers)

## 1. Anonymous Parameters

```
studid = 12345
cursor.execute(
    "SELECT name FROM Student WHERE sid=%s",
    (studid,) )
```

*This comma is no mistake, but needed with single parameters*

*parameter marker*

## 2. Named Parameters

```
studid = 12345
cursor.execute(
    "SELECT name FROM Student WHERE sid=%(sid)s",
    {'sid': studid} )
```

*named parameter marker*

### (3) Parameterized SQL & Host Variables

- Data transfer between DBMS and application
- Mapping of SQL domain types to data types of host language
- Python DB-API:
  - ▶ Host variables are normal *dynamically typed* Python variables; automatic conversion to/from SQL types done by psycopg in `execute()`:

```
studid = 12345
stmt = cursor.execute(
    "SELECT name FROM Student WHERE sid=%s",
    (studid,) )
```

- Note: in statement-level APIs such as ESQL/C:  
Host variables must be declared before usage

```
EXEC SQL BEGIN DECLARE SECTION;
    int studid = 12345;
    char sname[21];
EXEC SQL END DECLARE SECTION;
```

*Variables  
shared by host  
and SQL*

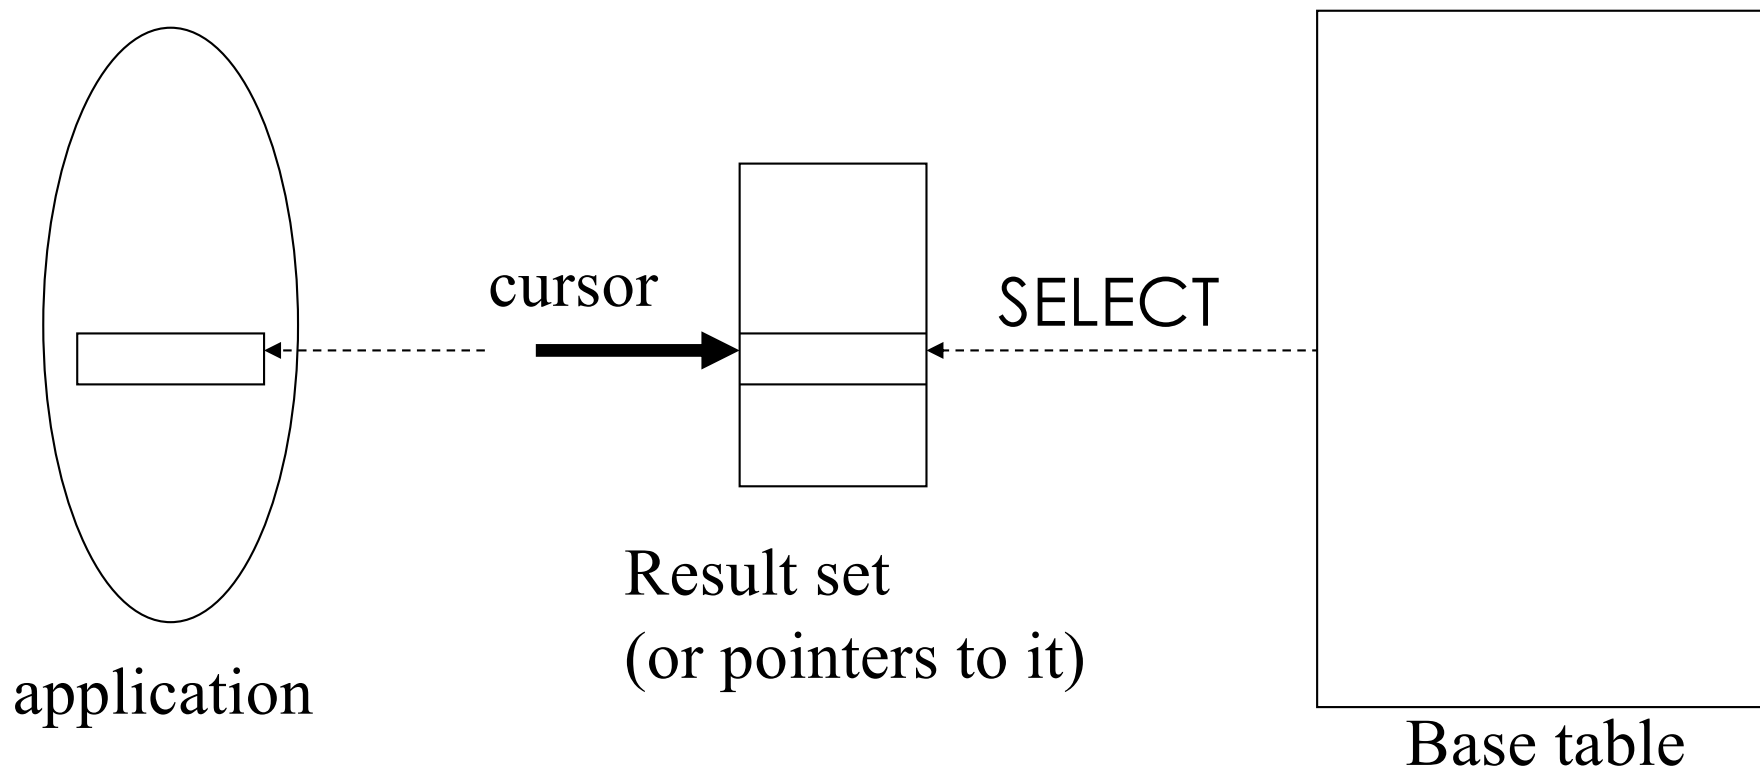
# Specifying Date/Time: Type Objects

- Providing date or time values is very database specific with different database configurations requiring particular formats
  - ▶ eg. "26.04.2016" (AU) versus "04/26/2014" (US)
- DB-API has helper objects to safely convert database types
  - ▶ **Date** ( *year* , *month* , *day* )
  - ▶ **Time** ( *hour* , *minute* , *second* )
  - ▶ **Timestamp** ( *year* , *month* , *day* , *hour* , *minute* , *second* )
  - ▶ **DateFromTicks** ( *ticks* )
  - ▶ **TimeFromTicks** ( *ticks* )
  - ▶ **TimestampFromTicks** ( *ticks* )
  - ▶ **Binary** ( *string* )

# (4) Buffer Mismatch Problem

## *(also: Impedance Mismatch)*

- **Problem:** SQL deals with tables (of arbitrary size); host language program deals with fixed size buffers
  - ▶ How is the application to allocate storage for the result of a SELECT statement?
- **Solution:** Cursor concept
  - ▶ Fetch a single row at a time



# Mapping of Sets: Cursor Concept

- **Result set** – set of rows produced by a SELECT statement
- **Cursor** – pointer to a row in the result set.
- Cursor operations:
  - ▶ *Declaration*
  - ▶ *Open* – execute SELECT to determine result set and initialize pointer
  - ▶ *Fetch* – advance pointer and retrieve next row (Python: fetchone() call)
  - ▶ *Close* – deallocate cursor

# Cursor in Python– via Cursor Interface

- Cursor concept with Python/psycopg:

```
cursor = conn.cursor()  
cursor.execute("SELECT title, name, address FROM Emp")  
row = cursor.fetchone()  
while row is not None:  
    print(row)  
    row = cursor.fetchone()  
cursor.close()
```

- Cursor objects are iterable, so shorter form is:

```
cursor = conn.cursor()  
cursor.execute("SELECT title, name, address FROM Emp")  
for row in cursor:  
    data = row[0] + "\t" + row[1] + "\t" + row[2] + "\n"  
    print(data)  
cursor.close()
```

You address result columns by position

# Cursor in Python – fetchAll()

- Fetchall() returns a Python list with all the result rows

- Good for small results

```
cursor.execute("SELECT title, name, address FROM Emp")  
resultset = cursor.fetchall()
```

```
cursor.close()
```

```
for row in resultset:  
    print(row)
```

- ▶ just be mindful that this will be **memory intensive** for large results



# Dictionary Cursors

- By default, psycopg returns tuples with `fetch()` / `fetchall()`
  - fields can only be addressed positionally
- As an extension, psycopg also supports dictionary cursors
  - Result is now a dictionary (associative array) which each field being named by the attribute names from the database schema

```
import psycopg2
from psycopg2.extras import RealDictCursor
...
curs = conn.cursor(cursor_factory=RealDictCursor)
curs.execute("SELECT title, name, address FROM Emp")
for row in curs:
    data = row['title'] + "\t" + row['name'] + "\n"
    print(data)
curs.close()
```

# NULL Handling in Python

- Remember: In SQL there is a special indication NULL used for unknown or inapplicable value of a column
  - ▶ Null value is not the same as 0 nor empty string

- In Python this shows as **None**:

```
cursor.execute("SELECT gender FROM Student ...")
result = cursor.fetchone()
if result[0] is None:
    # null value
else:
    # no null value
```

- Other languages require a special *indicator variable*. Eg. C:

```
EXEC SQL select gender into :gender:indicator
                                from Student where sid=4711;

if ( indicator == -1 )
{ /* null value */ }
else
{ /* no null value */ }
```

# Testing for Variable Exists / is None

- In Python, to check for existence of a variable versus whether it has None as value:

```
# Ensure variable is defined
```

```
try:
```

```
    x
```

```
except NameError:
```

```
    x = None
```

```
# Test whether variable is defined to be None
```

```
if x is None :
```

```
    some_fallback_operation()
```

```
else:
```


```
    some_operation(x)
```

[Source: <http://code.activestate.com/recipes/59892/> ]

## (5) Error Handling

- Multitude of potential problems
  - ▶ No database connection or connection timeout
  - ▶ Wrong login or missing privileges
  - ▶ SQL syntax errors
  - ▶ Empty results
  - ▶ NULL values
  - ▶ ...
- Hence always check database return values,
- Provide error handling code, resp. exception handlers
- Gracefully react to errors or empty results or NULL values
- **NEVER show database errors to end users**
  - ▶ Not only bad user experience, but huge security risk...

# You should avoid this!



Association for Computing Machinery  
Advancing Computing as a Science & Profession

ACM Order Rectification

The web site you are accessing has experienced an unexpected error.  
Please contact the website administrator.

The following information is meant for the website developer for debugging purposes.

Error Occurred While Processing Request

Element ORDERID is undefined in URL.

The error occurred in **D:\wwwroot\Public\rectifyCC\rectifyCC.cfm: line 463**

```
461 :      WHERE a.order_id = b.order_id
462 :      AND a.order_id = c.order_id
463 :      AND a.order_id = '#URL.orderID#'
464 :    </CFQUERY>
465 :
```

Resources:

- Check the [ColdFusion documentation](#) to verify that you are using the correct syntax.
- Search the [Knowledge Base](#) to find a solution to your problem.

Browser Mozilla/5.0 (Macintosh; U; Intel Mac OS X 10\_5\_8; en-us) AppleWebKit/531.9 (KHTML, like Gecko) Version/4.0.3 Safari/531.9

Remote Address 129.78.220.7

Referer

Cauta:

Also cf. error #... Of <http://www.sans.org/top25-software-errors/>

isys2120 sem2, 2022

# Error Handling with Python DB API

- Error handling via normal exception mechanism of Python
  - ▶ Errors and warnings are made available as Python exceptions
    - Warning                raised for warnings such as data truncation on insert, etc
    - Error                    exception raised for various db-related errors
- psycopg API extension:
  - ▶ Exception attributes for detailed SQL error codes and messages
  - ▶ **pgerror**    string of the error message returned by backend
  - ▶ **pgcode**    string with the **SQLSTATE** error code returned by backend

- Example:

```
try:
```

```
    psycopg2.connect(...)
```

```
except psycopg2.Error as e:
```

```
    print("Problem connecting to database:")
```

```
    print(e.pgerror)
```

```
    print(e.pgcode)
```

Demo purpose only. As said before:  
please do not directly print SQL exceptions ;)



# Exception Hierarchy of Python DB API

- The complete **Exception** inheritance hierarchy for the Python DB API is as follows:

```
StandardError
|
|_ Warning
|_ Error
    |_ InterfaceError
    |_ DatabaseError
        |_ DataError
        |_ OperationalError
            |_ psycopg2.extensions.QueryCancelError
            |_ psycopg2.extensions.TransactionRollbackError
        |_ IntegrityError
        |_ InternalError
        |_ ProgrammingError
        |_ NotSupportedError
```

[cf. <http://initd.org/psycopg/docs/module.html?highlight=connect#exceptions>]

# Reprise: Example of Python DB-API

```
import psycopg2
```

```
try:
```

```
# connect to the database  
conn = psycopg2.connect(database="postgres", user="test", password="secret")
```

```
# prepare to query the database  
curs = conn.cursor()
```

*cursor concept*

```
# execute a parameterised query
```

```
unit_of_study = "INFO2120"
```

```
curs.execute("""SELECT name  
                FROM Student NATURAL JOIN Enrolled  
                WHERE uos_code = %(uos)s""", {'uos': unit_of_study})
```

*dynamic SQL with  
safe parameter passing*

```
# loop through the resultset
```

```
for result in curs:
```

```
    print (" student: " + result[0])
```

```
# clean up
```

```
curs.close()
```

```
conn.close()
```

```
# error handling
```

```
except psycopg2.OperationalError as e:  
    print("unable to connect to database")
```

```
except Exception as e:  
    print("Error when querying database")  
    print(e)
```

*error handling*





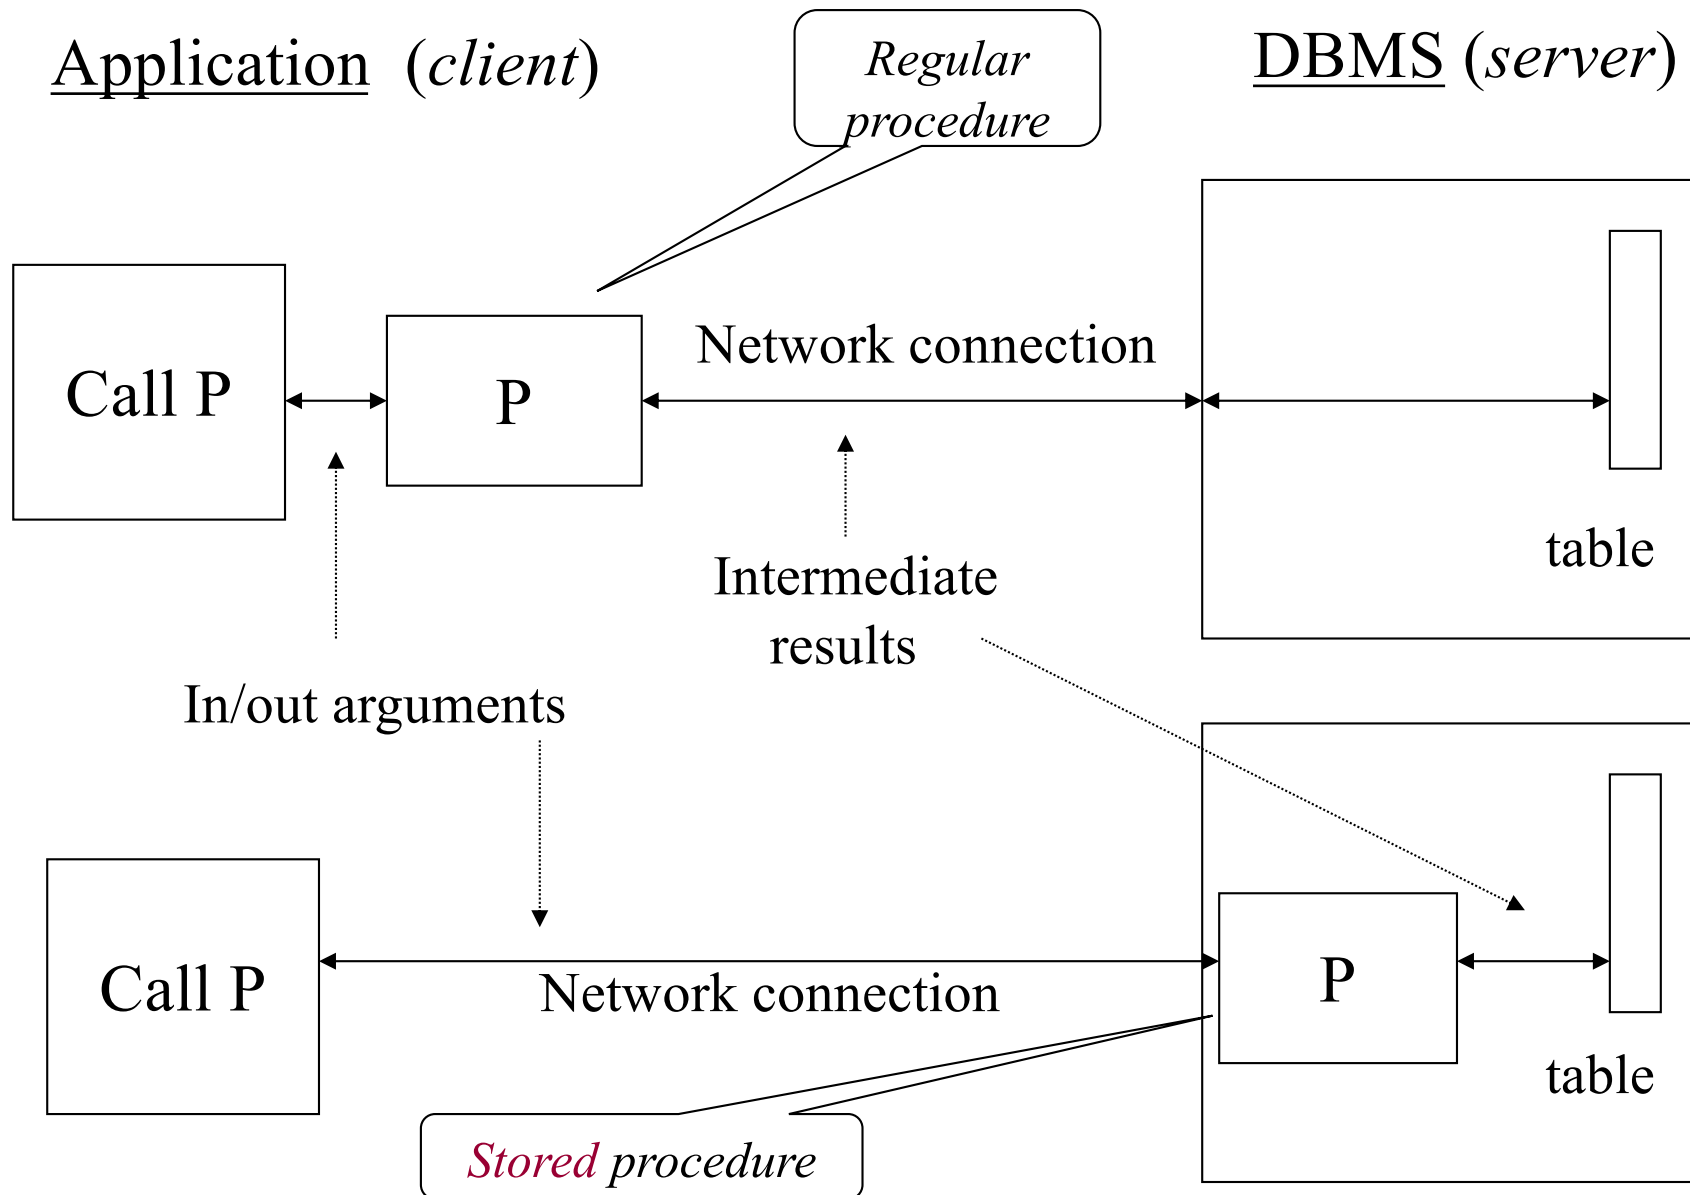
# Protecting a Web Database

- Be careful to check all parameters which can end up in such SQL statements!
  - ▶ Never trust user provided data!
- Use dynamic SQL statements with explicit, type-checked parameters (execute() function with parameter markers)
- Restrict the privileges of the user/role of the web application
  - ▶ E.g. with Oracle: Revoke EXECUTE privilege on Oracle standard packages when not needed. Specially for the PUBLIC role.
- Patch, patch, patch ;-)
- Also: NEVER directly return database error messages
  - ▶ Not very user-friendly AND it gives attackers hints

# Stored Procedures (Server-side application logic)

- Run application logic within the database server
  - ▶ Included as schema element (stored in DBMS)
  - ▶ Invoked by the application
- Advantages:
  - ▶ Central code-base for all applications
  - ▶ Improved maintainability
  - ▶ Additional abstraction layer  
(programmers do not need to know the schema)
  - ▶ Reduced data transfer
  - ▶ Less long-held locks
  - ▶ DBMS-centric security and consistent logging/auditing (important!)
- Note: although named procedures, can also be functions

# Stored Procedures



# SQL/PSM

- Stored Procedures not only have full access to SQL
- All major database systems provide extensions of SQL to a simple, general purpose language
  - ▶ SQL:1999 Standard: SQL/PSM
  - ▶ PostgreSQL: PL/pgSQL
  - ▶ Oracle: PL/SQL (syntax differs from PostgreSQL!!)
  - ▶ Microsoft SQL Server: T-SQL
- Extensions
  - ▶ Local variables, loops, if-then-else conditions
- Calling Stored Procedures: CALL statement
  - ▶ Example: **CALL *ShowNumberOfEnrolments()* ;**

# Procedure Declarations

## ■ Procedure Declarations (with SQL/PSM)

```
CREATE FUNCTION name ( parameter1, ..., parameterN ) AS  
    local variable declarations  
    procedure code;
```

## ■ Stored Procedures can have parameters

- ▶ of a valid SQL type (parameter types must match)
- ▶ three different modes
  - IN arguments to procedure
  - OUT return values
  - INOUT combination of IN and OUT

```
CREATE FUNCTION CountEnrolments ( IN uos VARCHAR ) AS  
    SELECT COUNT (*)  
        FROM Enrolled  
        WHERE uosCode = uos;  
  
CALL CountEnrolments ( 'INFO2120' );
```

# PostgreSQL: PL/pgSQL



(cf. <http://www.postgresql.org/docs/9.4/static/plpgsql.html>)

## ■ Extends SQL by programming language constructs

- ▶ Only knows functions! **CREATE FUNCTION** *name* **RETURNS** ... **AS**...
- ▶ Compound statements: **BEGIN** ... **END**;
- ▶ SQL variables: **DECLARE** *section*  
*variable-name* *sql-type*;
- ▶ Assignments: *variable* **:=** *expression*;
- ▶ IF statement: **IF** *condition* **THEN** ... **ELSE** ... **END IF**;
- ▶ Loop statements: **FOR** *var* **IN** *range* (**WHILE** *cond* )  
**LOOP** ... **END LOOP**;
- ▶ Return values: **RETURN** *expression*;
- ▶ Call statement: **CALL** *procedure(parameters)*;
- ▶ Transactions: **COMMIT**; **ROLLBACK**;



*Tip: CREATE OR REPLACE to avoid 'name-already-used'*

# PL/pgSQL Example

(cf. <http://www.postgresql.org/docs/9.4/static/plpgsql-structure.html>)

## ■ PL/pgSQL procedure declaration

**CREATE OR REPLACE FUNCTION**

*name* ( *parameter1*, ..., *parameterN* ) **RETURNS** *sqlType*

**AS \$\$**

**DECLARE**

*variable* *sqlType*;

...

**BEGIN**

...

**END;**

**\$\$ LANGUAGE plpgsql;**

*Tip: final delimiter must match the one used after AS*

*optional*

## ■ where parameterX is declared as (IN is default):

**[IN|OUT|IN OUT]** *name sqlType*

# PostgreSQL PL/pgSQL Example

```
CREATE OR REPLACE FUNCTION RateStudent
(studId INTEGER, uos VARCHAR) RETURNS CHAR AS $$
DECLARE
    grade CHAR;
    marks INTEGER;
BEGIN
    SELECT SUM(marks) INTO marks
    FROM Assessment
    WHERE sid=$1 AND uosCode=$2;
    IF ( marks>84 ) THEN grade := 'HD';
    ELIF ( marks>74 ) THEN grade := 'D';
    ELIF ( marks>64 ) THEN grade := 'CR';
    ELIF ( marks>50 ) THEN grade := 'P';
    ELSE
        grade := 'F';
    END IF;
    RAISE NOTICE 'Final grade is: %s', grade;
    RETURN grade;
END;
$$ LANGUAGE plpgsql;
```





# Python DB-API: Calling Stored Procedures

- **Cursor** objects have an explicit `callproc()` method
  - `cursor.callproc()` makes the OUT parameters available as resultset
- Example:

```
CREATE FUNCTION test(input VARCHAR,OUT output VARCHAR)
AS $$
BEGIN
    output := UPPER(input);
END $$ LANGUAGE plpgsql;
```

```
import psycopg2
conn = psycopg2.connect(...)
curs = conn.cursor()
input = "foo bar"
curs.callproc("test", [input] )
output = curs.fetchone()
print(output[0])
```

Pass all IN parameters  
as a list in order of the  
function declaration

OUT parameters are  
returned as resultset

# Language support of Stored Procedures

- Programming language virtual machine is often ‘integrated’ with DBMS
  - ▶ E.g. Java with Oracle
  - ▶ .Net CLR with IBM, Oracle, and SQL Server
  - ▶ PostgreSQL: Supports several scripting languages such as perl etc.
- But degree of integration differs heavily
  - ▶ If VM is in a different process from dbms, then performance often suffers

# Lessons Learned

## ■ Understand core issues for db-backed development

- ▶ Data and type conversion: *Host Variables*
- ▶ NULL value semantic: *Indicator variables* and testing methods
- ▶ Impedance Mismatch: *Cursor Concept*
- ▶ *Dynamic versus static SQL: security for parametrised queries*

## ■ Database APIs

- ▶ After lab08, you should in particular be able to work with small Python db programs

## ■ Server-side database programming

- ▶ How to use stored procedures to run code inside a DBMS
  - e.g. with PostgreSQL's pl/pgsql or with Oracle's PL/SQL
- ▶ Modern database engines provide virtual machine environments to run external code near the data

# References

- Kifer/Bernstein/Lewis (2nd edition)
  - ▶ Chapter 8
- Ramakrishnan/Gehrke (3rd edition - the 'Cow' book)
  - ▶ Chapter 6
- Ullman/Widom (3rd edition of 'First Course in Database Systems' )
  - ▶ Chapter 9 (*covers Stored Procedures, ESQL, CLI, JDBC and PHP*)
- Silberschatz/Korth/Sudarshan – Chapter 9

## Database Documentation:

- Python DB-API: <http://initd.org/psycopg/docs/>  
<https://wiki.python.org/moin/UsingDbApiWithPostgres>  
[http://www.tutorialspoint.com/postgresql/postgresql\\_python.htm](http://www.tutorialspoint.com/postgresql/postgresql_python.htm)
- The PostgreSQL Global Development Group: "PostgreSQL 8.2.4 Documentation", 2009.
- Oracle Corporation: "Oracle 10.1 Database Concepts", 2003.
- MySQL website: <http://www.mysql.com>